

如何参与 PostgreSQL 贡献

李建平

japinli@hotmail.com

2022.10.22

个人简介

- 成都文武信息技术有限公司
- 数据库内核开发
- PostgreSQL Contributor

公司简介

- 提供安全、可靠、经济、实用的高性能数据库和解决方案
- 专注于下一代高性能多模数据库
- 鸿鹄数据库
- 鸿鹄彩云超融合系统

目录

① 为什么参与贡献

② 如何参与贡献

③ 提交 Bugs

④ 代码审查

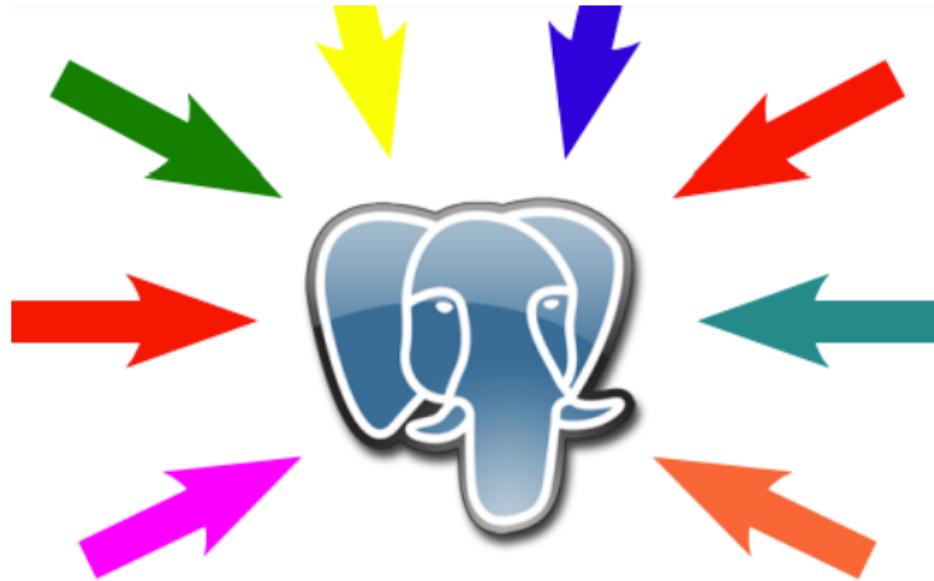
⑤ 代码贡献

为什么参与贡献

- ① 熟悉 PostgreSQL
- ② 了解内部原理
- ③ 自我提升



如何参与贡献



Josh Berkus 提供了 50 种方法

https://jberkus.github.io/pdf_presos/50_ways.pdf

常见的贡献方法

- 捐献
- 翻译
- 编写文档
- 提交 BUGs
- 代码审查
- 代码贡献

Donate

PostgreSQL Project assets

The PostgreSQL Community Association of Canada (PGCAC) is a special-purpose NPO setup at the request of the PostgreSQL Core Team to manage and protect key project assets such as domain names and trademarks.

[Donate to PGCAC](#)

General PostgreSQL efforts

PostgreSQL is an affiliated project of [Software in the Public Interest](#). Funds donated to PostgreSQL are used to sponsor general PostgreSQL efforts. These funds are managed by the Fund raising group.

[Donate to PostgreSQL](#)

PostgreSQL Europe (PgEU)

PostgreSQL Europe is a European region non profit focusing on the support of other PostgreSQL groups in Europe.

[Donate to PgEU](#)

PostgreSQL United States (PgUS)

PostgreSQL United States is a non profit for PostgreSQL in the United States. It focuses on Education, User Groups, and Advocacy.

<https://www.postgresql.org/about/donate/>

常见的贡献方法

- 捐献
- 翻译
- 编写文档
- 提交 BUGs
- 代码审查
- 代码贡献

Servers

The servers that power the services of www.postgresql.org are provided by different companies and organisations around the world.

Name	Provider	Dedicated	Specs	OS	Location	Usage
aws	Amazon Web Services	No	Dynamic	Proprietary	Global	Limited VM hosting
baku	conova communications GmbH	Yes	6 core, 32Gb	Debian GNU/Linux 10	Salzburg, Austria	internal use
barzan	conova communications GmbH	Yes	8 core, 64Gb	Debian GNU/Linux 10	Salzburg, Austria	internal use
poseidon	EDB	Yes	8 core, 384GB	CentOS release 7.x	Wantage, Oxfordshire, UK	RPM build server
postgresql02	EDB	Yes	4 core, 8GB	Red Hat Enterprise Linux Server	St. Louis, Missouri, United States	Virtual machine host, independent of the pgInfras platform
ranza	Rackspace	Yes	12 core, 32Gb	Debian GNU/Linux 10	United States	Virtual machine host
regula	Redpill Linpro	Yes	48 core, 256Gb	Debian GNU/Linux 10	Oslo, Norway	Virtual machine host
relva	Rackspace	Yes	12 core, 32Gb	Debian GNU/Linux 10	United States	Virtual machine host
rigel	Rackspace	Yes	12 core, 32Gb	Debian GNU/Linux 10	United States	Virtual machine host
risa	Rackspace	Yes	12 core, 32Gb	Debian GNU/Linux 10	United States	Virtual machine host
thallion	conova communications GmbH	Yes	32 core, 256Gb	Debian GNU/Linux 10	Salzburg, Austria	Virtual machine host
triona	conova communications GmbH	Yes	16 core, 128Gb	Debian GNU/Linux 10	Salzburg, Austria	Virtual machine host

<https://www.postgresql.org/about/servers/>

常见的贡献方法

● 捐献

● 翻译

● 编写文档

● 提交 BUGs

● 代码审查

● 代码贡献

NLS Status Tables

The tables below shows which areas of the PostgreSQL source are prepared for translation, and the progress of those translations. The numbers are percents of translated messages. By following the links, you can download freshly baked PO files that are merged up against the latest program sources. No PostgreSQL source tree is needed for translation work.

The highlighted boxes are the files that have a sufficient fraction of strings translated to be considered for release. Files that are 100% translated are specially highlighted. All other files will not be part of the PostgreSQL release.

The grey numbers are translations that do not exist yet, but where this web site has initialized the translations files for you with data from the existing translations for the same language.

You can work in several branches. Please see the status information below and follow the announcements on the [mailing list](#) ([subscribe](#)) about when a branch is about to be released.

- [15 branch](#) — complete this now
- [14 branch](#) — maintenance
- [13 branch](#) — maintenance
- [12 branch](#) — maintenance
- [11 branch](#) — maintenance
- [10 branch](#) — maintenance

Other information: [contribution/translator information \(wiki\)](#)

15 branch

Last update: Thu Oct 13 2022 08:22:03 GMT+0800 (China Standard Time)

	af	cs	de	el	es	fa	fr	he	hr	hu	id	it	ja	ka	ko	nb	nl	pl	pt_BR	xo	xw	sk	sl	sv	ta	tr	uk	vi	zh_CN	zh_TW	30
expg (331 strings)	32	30	32	30		30	3	3	1	10	30	30	30	30	2	1	32	30	8	30	11	2	30	2	30	30	30	30	30	18 (95)	
expelb (32 strings)	30	30	30	30		30					30	30	30	30	30			30	30		30			30	30	30	30	30	30	17 (100)	
initdb (188 strings)	72	30	81	30		30	37	5	1	17	38	30	30	30	79	5	6	37	30	24	30	14	18	30	20	25	98	38	81	24	23 (67)
libpq (265 strings)	19	83	30	92	30		30	63	17	1	3	22	30	30	85	16	60	60	1	30	19	27	30	38	23	99	1	92	46	24 (69)	
pg_amcheck (97 strings)	20	30	90	30	1	30	12	2	1	4	15	30	30	22	5	15	28	8	30	9	7	30	7	19	30	4	90	11	10 (98)		
pg_archivecleanup (27 strings)	33	30	33	30		30	7			7	22	30	30	88	11	3	37	30	11	30	14	14	30	7	31	30	37	38	14	16 (87)	
pg_basebackup (331 strings)	24	30	24	30		30	15	1	1	6	20	30	30	24	3	2	16	14	2	30	6	5	30	1	20	98	20	74	7	19 (65)	
pg_checksums (39 strings)	22	30	29	30		30	8	3		10	20	30	30	22	6	1	20	15	30	10	10	30	6	66	30	6	29	16	14 (91)		
pg_config (43 strings)	22	30	30	30	2	30	83	6	4	18	83	30	30	82	69	9	83	30	29	30	11	32	30	22	83	30	81	30	23 (88)		
pg_controldata (92 strings)	1	30	30	98	30	21	30	8	3	20	8	32	30	98	21	5	72	30	44	30	23	25	30	34	83	30	78	98	44	25 (73)	
pg_ctl (1 string)	30	30	98	30	1	30	78	1	1	8	82	30	30	98	6	5	72	30	27	30	3	26	30	45	82	30	15	98	56	22 (80)	

<https://babel.postgresql.org>

常见的贡献方法

- 捐献
- 翻译
- 编写文档
- 提交 BUGs
- 代码审查
- 代码贡献



常见的贡献方法

- 捐献
- 翻译
- 编写文档
- 提交 BUGs
- 代码审查
- 代码贡献

Submit Bug Report

Please ensure you have read the [bug reporting guidelines](#) before reporting a bug. In particular, please re-read the [documentation](#) to verify that what you are trying is possible. If the documentation is not clear, please report that, too; it is a documentation bug. If a program does something different from what the documentation says, that is also a bug.

Poor performance is not necessarily a bug. Read the documentation or ask on one of the [mailing lists](#) for help in tuning your applications. Failing to comply to the SQL standard is not necessarily a bug either, unless compliance for the specific feature is explicitly claimed.

Before you continue, check on the [TODO list](#) and in the [FAQ](#) to see if your bug is already known. If you cannot decode the information on the TODO list, report your problem so we can clarify the TODO list.

If you believe you have found a security issue, please send an email to security@postgresql.org. All other bugs will be forwarded to the [pgsql-bugs](#) mailing list where they will be publicly archived.

Make sure you are running the latest available minor release for your major [version](#) before reporting a bug. The current list of supported versions is 14.5, 13.8, 12.12, 11.17, 10.22.

This bug report form should only be used for reporting bugs and problems with the PostgreSQL database. Problems with database connectors such as ODBC and JDBC, graphical administration tools such as pgAdmin or other external projects [should not be reported here](#); please report to those projects directly. For products closely connected with PostgreSQL, there may be an appropriate [mailing list](#) available.

By submitting a bug through this form you agree that all contents of the bug report, including your personal information as listed, will be posted to the [public](#) [pgsql-bugs](#) mailing list, and archived in the [public](#) list archives.

Name:

Japin Li

Email:

japinli@hotmail.com

PostgreSQL version:

-- Select version

Operating system:

Short description:

<https://www.postgresql.org/account/submitbug/>

常见的贡献方法

- 捐献
- 翻译
- 编写文档
- 提交 BUGs
- 代码审查
- 代码贡献

Commitfest 2022-11

Search/filter Shortcuts + New patch

Status summary: Needs review: 131, Waiting on Author: 68, Ready for Committer: 23, Committed: 27, Returned with Feedback: 2, Withdrawn: 4, Rejected: 1, Total: 256.

Active patches

Patch	Status	Ver	Author	Reviewers	Committer	Num cfs	Latest activity	Latest mail
Bug Fixes								
pgbench: using prepared BEGIN statement in a pipeline could cause an error	Needs review	15	Yugo Nagata (yugo.nagata)	Kyotaro Horiguchi (horiguti), Fabien Coelho (calvini), Daniel Gustafsson (dege)		7	2022-10-12 05:10	2022-09-30 01:07
Error "initial slot snapshot too large" in create replication slot	Waiting on Author		Kyotaro Horiguchi (horiguti), Dilip Kumar (dilip-kumar)			5	2022-10-12 05:10	2022-10-12 05:10
Fix checkpoint sync request queue problems	Waiting on Author		Thomas Munro (maddice)			3	2022-10-12 05:11	2022-03-16 03:04
Nonreplayable XLog records by means of overflows and >MaxAllocSize lengths	Waiting on Author	status	Matthias van de Meent (mmeent)	Michael Paquier (michael-kun)	michael-kun	3	2022-10-05 07:46	2022-10-05 07:46
Remove pgwin32_ls_junction(), add lstat() for Windows	Waiting on Author		Thomas Munro (maddice)			3	2022-10-12 05:11	2022-10-12 04:05
Fix dsa_free() to re-bin segment	Needs review		Dongming Liu (mferyes)			3	2022-10-12 05:07	2022-04-06 07:10
Fix recovery conflict SIGUSR1 handling	Needs review		Thomas Munro (maddice)	Michael Paquier (michael-kun)		3	2022-10-06 06:41	2022-07-26 23:22
Data is copied twice when specifying both child and parent table in publication	Waiting on Author		wei wang (wangwei), Yu Shi (shiyifinst)	Amrit Kapila (amrikapila), Takamichi Osumi (tataapppi), Hou Zhijie (houzj)		3	2022-10-12 05:11	2022-10-05 15:04
Fix assertion failure with barriers in parallel hash join	Ready for Committer	status	Thomas Munro (maddice), Melanie Plagman (melanieplagman)	David Geier (geidav)		3	2022-06-06 05:11	2022-06-06 15:01
pg_rewind: warn when checkpoint hasn't happened after promotion	Needs review		James Coleman (jcole1989)	Yogesh Sharma (sharmay)		3	2022-10-12 05:12	2022-07-06 02:38

<https://commitfest.postgresql.org/40/>

常见的贡献方法

- 捐献
- 翻译
- 编写文档
- 提交 BUGs
- 代码审查
- 代码贡献

PostgreSQL BuildFarm Status

Shown here is the latest status of each farm member for each branch it has reported on in the last 30 days.

Use the farm member link for history of that member on the relevant branch.

Legend: = cassandra = debug = gssapi = krb5 = lvm = nis = openssl
 = pam = perl = python = tap-tests = tcl = thread-safety = xpath = xml

Alias	System	Status	Flags
vulpes	fedora 27 gcc 7.3.1 ppc64le	00:42 ago OK [56c19fe] Config	
morepork	OpenBSD OpenBSD 6.9 clang clang 10.0.1 x64	00:50 ago OK [56c19fe] Config	
conchuela	DragonFly BSD DragonFly BSD 6.2.2 gcc 8.3 x86_64	01:00 ago OK [56c19fe] Config	
flaviventeris	Debian Sid gcc snapshot x86_64	01:05 ago OK [56c19fe] Config	
loach	FreeBSD FreeBSD 12.2 clang clang 10.0.1 x86_64	01:11 ago OK [56c19fe] Config	
serinus	Debian Sid gcc snapshot x86_64	01:11 ago OK [56c19fe] Config	
gualibaseurus	Debian GNU/Linux 11 (bullseye) gcc 10.2.1 x86_64	01:16 ago OK [56c19fe] Config	
curculio	OpenBSD 5.9 gcc 4.2.1 x86_64	01:20 ago OK [56c19fe] Config	
sidewinder	NetBSD NetBSD 9.2 gcc clang 12.0.1 x86_64	01:30 ago OK [56c19fe] Config	
butterflyfish	Photon 2.0 Gcc 6.3.0 x86_64	01:36 ago OK [56c19fe] Config	
myna	Photon 3.0 Gcc 6.3.0 x86_64	01:36 ago OK [56c19fe] Config	
clam	RHEL 7.1 IBM Advance Toolchain GCC 5.2.1 (Advance-Toolchain-at9.0) ppc64le	01:36 ago OK [56c19fe] Config	
gombessa	OpenBSD 6.8 Clang 10.0.1 amd64	01:37 ago OK [56c19fe] Config	
rhinoceros	CentOS Linux 7.1 gcc 4.8.3 20140911 (Red Hat 4.8.3-9) x86_64	01:44 ago OK [56c19fe] Config	
plover	OpenBSD 6.9 clang 10.0.1 amd64	01:47 ago OK [56c19fe] Config	
warbler	FreeBSD 14 Current Clang 13.0.0 x86_64	02:02 ago OK [56c19fe] Config	
jay	openSUSE openSUSE 15.3 clang clang 11 x86_64	02:03 ago OK [56c19fe] Config	
mule	Debian 11.2 gcc 10.2.1 amd64	02:06 ago OK [56c19fe] Config	
peripatus	FreeBSD HEAD clang 14.0.5 x86_64	02:16 ago OK [56c19fe] Config	

<https://buildfarm.postgresql.org>

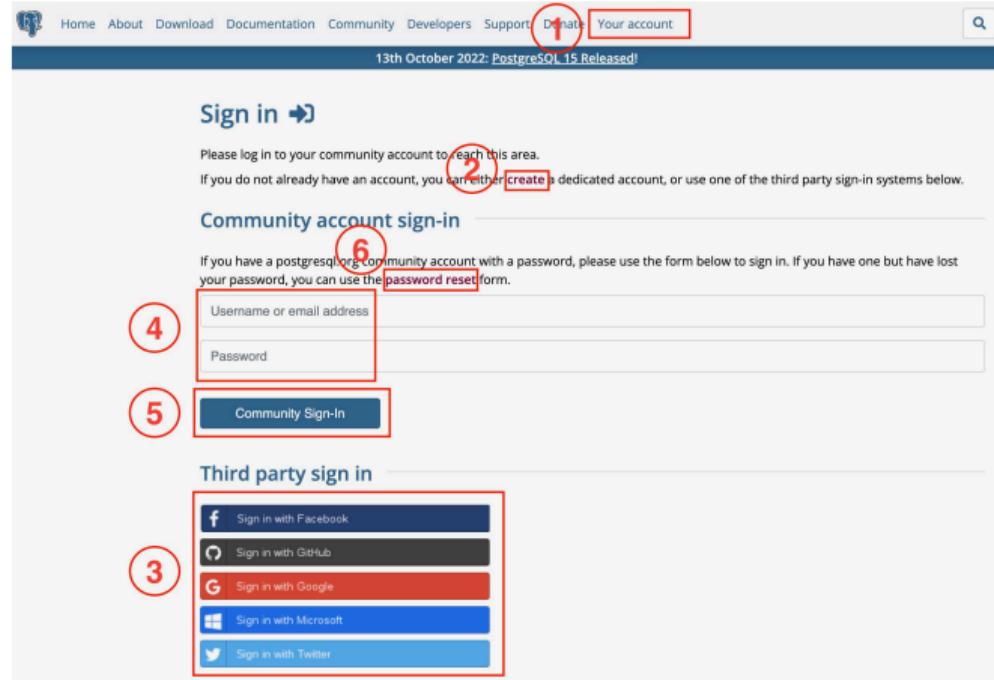
准备工作

- PostgreSQL 账号

- 账号管理

- 订阅邮件列表

- 邮件列表说明



<https://www.postgresql.org/account/>

准备工作

- PostgreSQL 账号

- 账号管理

- 订阅邮件列表

- 邮件列表说明

New Account

To sign up for a free community account, enter your preferred userid and email address. Note that a community account is only needed if you want to submit information - all content is available for reading without an account. A confirmation email will be sent to the specified address, and once confirmed a password for the new account can be specified.

Username:

First name:

Last name:

Email:

Repeat email:

Captcha:

Sign up

<https://www.postgresql.org/account/signup/>

准备工作

- PostgreSQL 账号

- 账号管理

- 订阅邮件列表

- 邮件列表说明

The screenshot shows the 'Edit User Profile' page. A red box labeled '1' highlights the 'Change password' link in the 'Quick Links' sidebar. A red box labeled '2' highlights the 'Primary Email Address' field containing 'japinli@hotmail.com'. A red box labeled '3' highlights the 'First Name' field containing 'Japin'. A red box labeled '4' highlights the large 'Secondary email addresses' input area. A red box labeled '5' highlights the 'New email' input field in the 'Add new email address' section.

Quick Links

- Your account
- Profile
- Mailing Lists
- Subscriptions
- Submitted data
- Events
- Products
- Professional Services
- Organisations
- Change password
- Logout

Edit User Profile

Username: **japin**

Primary Email Address
To change the primary email address,
first add it as a secondary address
below

First Name: **Japin**

Last Name: **Li**

SSH key
Paste one or more public keys in
OpenSSH format, one per line.

Block OAuth login
Disallow login to this account using
OAuth providers like Google or
Microsoft.

Secondary email addresses

You can add one or more secondary email addresses to your account, which can be used for example to subscribe to mailing lists.

Add new email address

New email: **(5)**

Repeat email:

Save

<https://www.postgresql.org/account/profile/>

准备工作

- PostgreSQL 账号

- 账号管理

- 订阅邮件列表

- 邮件列表说明

The screenshot shows the PostgreSQL mailing list management interface at <https://lists.postgresql.org/manage/>. It includes sections for 'Your subscriptions', 'Subscribe', 'Manage email addresses', and 'Global configuration'.

Your subscriptions: Shows current subscriptions to three lists: pgsql-advocacy@lists.postgresql.org, pgsql-bugs@lists.postgresql.org, and pgsql-hackers@lists.postgresql.org. Each entry has 'Mail delivery' status (Enabled) and 'Actions' buttons for View archive, Edit subscription, Unsubscribe, and Send test mail.

Subscribe: A form to add a new mailing list subscription. It includes fields for 'List' (pgsql-admin@pgsql-admin) and 'Address' (japinli@hotmail.com), followed by a 'Subscribe' button.

Manage email addresses: A table showing one email address: japinli@hotmail.com, which is 'Confirmed'. There is a 'Block' button next to it.

Global configuration: A section where configuration settings apply to all list subscriptions. It includes a checkbox for 'Don't receive an extra copy of mails when listed in To or CC fields' and a 'Save' button.

Log out: A link to log out of the account.

<https://lists.postgresql.org/manage/>

准备工作

- PostgreSQL 账号
 - pgsql-general
 - pgsql-performance
 - pgsql-sql
- 账号管理
 - pgsql-bugs
 - pgsql-docs
 - pgsql-hackers
 - pgsql-translators
- 订阅邮件列表
 - pgsql-zh-general
- 邮件列表说明

<https://www.postgresql.org/list/>

提交 Bugs

● 页面提交（推荐）

● 安全相关 Bugs

● 邮件列表

Submit Bug Report

Please ensure you have read the [bug reporting guidelines](#) before reporting a bug. In particular, please re-read the [documentation](#) to verify that what you are trying is possible. If the documentation is not clear, please report that, too; it is a documentation bug. If a program does something different from what the documentation says, that is also a bug.

Poor performance is not necessarily a bug. Read the documentation or ask on one of the [mailing lists](#) for help in tuning your applications. Failing to comply to the SQL standard is not necessarily a bug either, unless compliance for the specific feature is explicitly claimed.

Before you continue, check on the [TODO list](#) and in the [FAQ](#) to see if your bug is already known. If you cannot decode the information on the TODO list, report your problem so we can clarify the TODO list.

If you believe you have found a security issue, please send an email to security@postgresql.org. All other bugs will be forwarded to the [public pgsql-bugs](#) mailing list where they will be publicly archived.

Make sure you are running the latest available minor release for your major [version](#) before reporting a bug. The current list of supported versions is 14.5, 13.8, 12.12, 11.17, 10.22.

This bug report form should only be used for reporting bugs and problems with the PostgreSQL database. Problems with database connectors such as ODBC and JDBC, graphical administration tools such as pgAdmin or other external projects [should not be reported here](#); please report to those projects directly. For products closely connected with PostgreSQL, there may be an appropriate [mailing list](#) available.

By submitting a bug through this form you agree that all contents of the bug report, including your personal information as listed, will be posted to the [public pgsql-bugs](#) mailing list, and archived in the [public](#) list archives.

Name:

japin Li

Email:

japinli@hotmail.com

PostgreSQL version:

** Select version

Operating system:

Short description:

提交 Bugs

- 页面提交（推荐）

Submit Bug Report

Please ensure you have read the [bug reporting guidelines](#) before reporting a bug. In particular, please re-read the [documentation](#) to verify that what you are trying is possible. If the documentation is not clear, please report that, too; it is a documentation bug. If a program does something different from what the documentation says, that is also a bug.

Poor performance is not necessarily a bug. Read the documentation or ask on one of the [mailing lists](#) for help in tuning your applications. Failing to comply to the SQL standard is not necessarily a bug either, unless compliance for the specific feature is explicitly claimed.

Before you continue, check on the [TODO list](#) and in the [FAQ](#) to see if your bug is already known. If you cannot decode the information on the TODO list, report your problem so we can clarify the TODO list.

If you believe you have found a security issue, please send an email to security@postgresql.org. All other bugs will be forwarded to the [pgsql-bugs](#) mailing list where they will be publicly archived.

Make sure you are running the latest available minor release for your major [version](#) before reporting a bug. The current list of supported versions is 14.5, 13.8, 12.12, 11.17, 10.22.

This bug report form should only be used for reporting bugs and problems with the PostgreSQL database. Problems with database connectors such as ODBC and JDBC, graphical administration tools such as pgAdmin or other external projects [should not be reported here](#); please report to those projects directly. For products closely connected with PostgreSQL, there may be an appropriate [mailing list](#) available.

By submitting a bug through this form you agree that all contents of the bug report, including your personal information as listed, will be posted to the [public pgsql-bugs](#) mailing list, and archived in the [public](#) list archives.

Name:

Japin Li

Email:

japinli@hotmail.com

PostgreSQL version:

** Select version

Operating system:

Short description:

- 安全相关 Bugs

- 邮件列表

提交 Bugs

- 页面提交（推荐）

- 安全相关 Bugs

- 邮件列表

Re: Added schema level support for publication.

From: Amit Kapila <amit@dkoskapila16@gmail.com>
To: Li Japin <japinli@hotmail.com>
Cc: Bharath.Rupireddy <bharathdotrupireddyforpostgres@gmail.com>, vignesh C <vignesh71@gmail.com>, PostgreSQL Hackers <pgsql-hackers@list(dot)postgresql(dot)org>
Subject: Re: Added schema level support for publication.
Date: 2021-01-11 11:57:22
Message-ID: CAAMjACVw0Fqz2ScrgbpuW9s0Htg1g@2a.430U/bgrU.wx@mail.gmail.com
Views: Raw Message | Whole Thread | Download mbox | Resend email
Thread: 2021-01-11 11:57:22 from Amit Kapila <amit@dkoskapila16@gmail.com>
Lists: pgsql-hackers

On Mon, Jan 11, 2021 at 4:29 PM Li Japin <japinli@hotmail.com> wrote:
>
> On Mon, Jan 11, 2021, at 5:04 PM, Bharath Rupireddy <bharath(dot)rupireddyforpostgres(at)gmail(dot)com> wrote:
>
> On Mon, Jan 11, 2021 at 1:29 PM jatin <jatinlist@hotmail(dot)com> wrote:
>
> Sorry, I misunderstood. After the test (see #11530), I found that if we do not insert data
> between step (5) and (6), it will not skip the new records, however, if we insert
> data between step (5) and (6), it will skip the new records.
>
>
> It might be a bug.
>
Can you check pg_publication_rel and pg_subscription_rel? Also, this
is not related to the feature proposed in this thread, so it is better
to start a new thread to conclude whether this is a bug or not.

--
With regards,
Amit Kapila

Logical Replication - behavior of ALTER PUBLICATION .. DROP TABLE and ALTER SUBSCRIPTION .. REFRESH PUBLICATION

From: Bharath.Rupireddy <bharathdotrupireddyforpostgres@gmail.com>
To: PostgreSQL Hackers <pgsql-hackers@list(dot)postgresql(dot)org>
Subject: Logical Replication - behavior of ALTER PUBLICATION .. DROP TABLE and ALTER SUBSCRIPTION .. REFRESH PUBLICATION
Date: 2021-01-11 13:21:15
Message-ID: CAJjACVw0Fqz2ScrgbpuW9s0Htg1g@2a.430U/bgrU.wx@mail.gmail.com
Views: Raw Message | Whole Thread | Download mbox | Resend email
Thread: 2021-01-11 13:21:15 from Bharath.Rupireddy <bharathdotrupireddyforpostgres(at)gmail(dot)com>
Lists: pgsql-hackers

Hi,

While providing thoughts on the design in [1], I found a strange
behaviour with the subject. The use case is shown below as a sequence
of steps that need to be run on publisher and subscriber to arrive at
the strange behaviour. In step 5, the table is dropped from the
publisher. Now in step 6, the publisher runs a refresh on the subscriber.
From here onwards, the expectation is that no further
inserts into the publisher table have to be replicated on to the
subscriber, but the opposite happens i.e. the inserts are still
replicated to the subscriber. ISTM as a bug. Let me know if I'm
missing anything.

Thoughts?

代码审查

- 审查流程

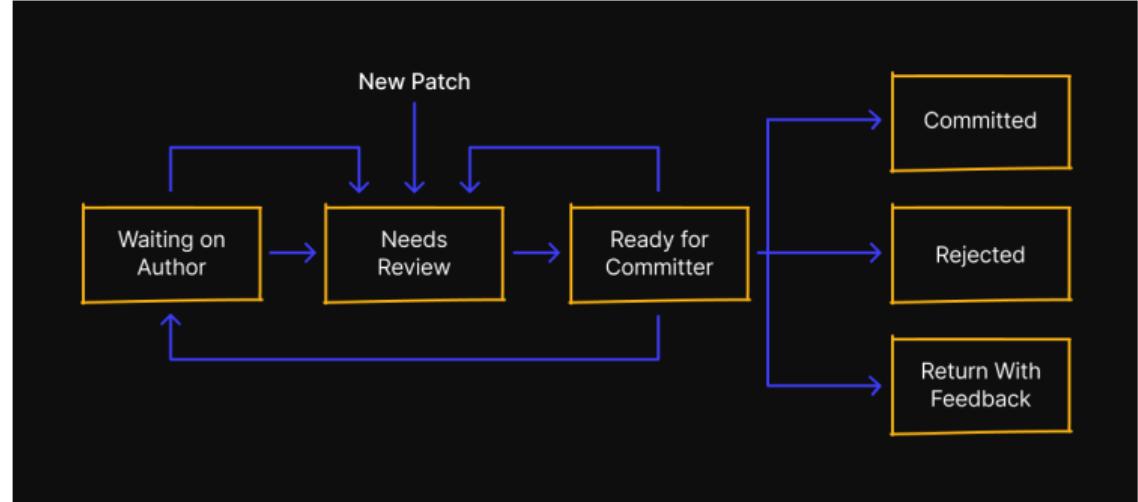
- 查找功能

- 代码审查

- 文档

- 功能

- 测试反馈



图片源自 State of PostgreSQL: How to Contribute to PostgreSQL and the Community

代码审查

● 审查流程

● 查找功能

● 代码审查

• 文档

• 功能

• 测试反馈

Commitfest 2022-11

Search/filter Shortcuts New patch

Status summary: Needs review: 129, Waiting on Author: 70, Ready for Committer: 22, Committed: 22, Withdrawn: 4, Rejected: 1, Returned with Feedback: 2, Total: 251.

Active patches

Patch	Status	Ver	Author	Reviewers	Committer	Num chs	Lastest activity	Lastest mail
Bug Fixes								
pgbench: using prepared BEGIN statement in a pipeline could cause an error	Needs review		Yugo Nagata (yugo.nagata)	Kyotaro Honjuchi (honjuki), Fabien Coelho (calvin), Daniel Gustafsson (tege)		7	2022-10-19 05:19	2022-09-30 01:07
Error "Initial slot snapshot too large" in create replication slot	Working on Author		Kyotaro Honjuchi (honjuki), Dilip Kumar (dilip.kumar)			5	2022-10-12 05:10	2022-10-12 05:10
Fix checkpoint sync request queue problems	Working on Author		Thomas Munro (macdot)			3	2022-10-12 05:11	2022-03-16 02:04
Nonplayable XLog records by means of overflow and >MaxAllocSize lengths	Working on Author		Mathias van de Meent (mvanmeent)	Michael Paquier (michael-kun)	michaell-kun	3	2022-10-05 07:46	2022-10-05 07:46
Remove pgwini32_ie_junction(), add list() for Windows	Working on Author		Thomas Munro (macdot)			3	2022-10-12 05:11	2022-10-12 04:05
Fix das_head to n-bin segment	Needs review		Dongming Liu (lirongye)			3	2022-10-12 05:07	2022-04-06 07:10
Fix recovery conflict SIGUSR1 handling	Needs review		Thomas Munro (macdot)	Michael Paquier (michael-kun)		3	2022-10-06 06:41	2022-07-26 23:22
Data is copied twice when specifying both child and parent table in publication	Working on Author		wei wang (wangwei), Yu Shi (shiyi,frst)	Amrit Kapila (amritkapila), Takamichi Osurni (takatapp), Hou Zhihe (houz)		3	2022-10-12 05:11	2022-10-05 15:04
Fix assertion failure with barriers in parallel hash join	Ready for Committer		Thomas Munro (macdot), Melanie Plagman (melanieplagman)	David Geer (geede)		3	2022-10-12 05:11	2022-06-06 15:01
pg_rewind: warn when checkpoint hasn't happened after promotion	Needs review		James Coleman (cole1988)	Yogesh Sharma (jharmay)		3	2022-10-12 05:12	2022-07-06 02:38
pg_visibility & pg_check_visibility yields false positive when working in parallel with autovacuum	Working on Author		Daniel Shlepkov (dshelepkov)			3	2022-10-12 05:15	2022-10-12 08:14
pg_stat_activity: avoid showing stale-active with wait_event=ClientRead	Working on Author		Yury Sokolov (funny, talcor), Mikhail Zhulin (mizuka-al-pgsql)	Justin Pryby (justinprby), Kamal Priya (kamalp)		2	2022-10-12 07:48	2022-07-09 02:02

代码审查

● 审查流程

● 查找功能

● 代码审查

• 文档

• 功能

• 测试反馈

pgbench: using prepared BEGIN statement in a pipeline could cause an error

	Edit Comment/Review Change Status
Title	pgbench: using prepared BEGIN statement in a pipeline could cause an error
Topic	Bug Fixes
Created	2021-07-18 14:35:27
Last modified	2022-10-12 05:10:52 (1 day, 22 hours ago)
Latest email	2022-09-30 01:07:43 (2 weeks ago)
Status	2022-11: Review needed 2022-09: Moved to next CF 2022-07: Moved to next CF 2022-03: Moved to next CF 2022-01: Moved to next CF 2021-11: Moved to next CF 2021-09: Moved to next CF
Target version	13
Authors	Yugo Nagata (yugo.nagata)
Reviewers	Kyotaro Horiguchi (horugi), Fabien Coelho (calvin), Daniel Gustafsson (drg)
Committer	
Links	
Emails	pgbench: using prepared BEGIN statement in a pipeline could cause an error <small>(x)</small> First at 2021-07-18 05:00:12 by Yugo NAGATA <yngtai@arccos.co.jp> Latest at 2022-09-30 01:07:43 by Yugo NAGATA <yngtai@arccos.co.jp> Latest attachment yb-pgbench-prepare.patch at 2022-09-30 01:07:43 from Yugo NAGATA <yngtai@arccos.co.jp> + Attach thread
History	2021-11-04 13:11:37 Daniel Gustafsson (dt.gustafsson) Changed targetversion to 13 2021-10-04 07:31:24 Jaime Casanova (jcasanov) Closed in commitfest 2021-09 with status: Moved to next CF 2021-07-18 14:35:44 Yugo Nagata (yugo.nagata) Changed authors to Yugo Nagata (yugo.nagata)

代码审查

● 审查流程

● 查找功能

● 代码审查

• 文档

• 功能

• 测试反馈

pgbench: using prepared BEGIN statement in a pipeline could cause an error

From: Yugo NAGATA <nagata(at)sraoss(dot)co(dot)jp>
To: pgsql-hackers(at)postgresql(dot)org
Subject: pgbench: using prepared BEGIN statement in a pipeline could cause an error
Date: 2021-07-16 06:30:13

Message-ID: 20210716153013.fc53b1c780bd6fccc07a7f0d1@sraoss.co.jp

Views: Raw Message | Who's Thread | Download mbox | Resend email

Thread: 2021-07-16 06:30:13 from Yugo NAGATA <nagata(at)sraoss(dot)co(dot)jp>
Lists: pgsql-hackers

Hello, [查看所有邮件](#) [选择查看指定发件人的邮件](#)

I found that using "BEGIN ISOLATING LEVEL SERIALIZABLE" in a pipeline with prepared statements makes pgbench abort.

```
$ cat pipeline.sql
\startpipeline
begin isolation level repeatable read;
select 1;
\end;
\endpipeline

$ pgbench -f pipeline.sql -N prepared -t 1
pgbench (15devrel)
starting vacuum...end.
pgbench: error: client 0 script 0 aborted in command 4 query 0:
transaction type: pipeline.sql
scaling factor: 1
query mode: prepared
number of clients: 1
number of threads: 1
number of transactions per client: 1
number of transactions actually processed: 0/1
pgbench fatal: Run was aborted; the above results are incomplete.
```

The error that occurred in the backend was
"ERROR: SET TRANSACTION ISOLATION LEVEL must be called before any query".

After investigating this, now I've got the cause as below.

1. The commands in the script are executed in the order. First, pipeline mode starts at \startpipeline.
2. Parse messages for all SQL commands in the script are sent to the backend because it is first time to execute them.
3. An implicit transaction starts, and this is not committed yet because sync message is not sent at that time in pipeline mode.
4. All prepared statements are sent to the backend.
5. After processing \endpipeline, sync is issued and all sent commands are executed.
6. However, the BEGIN doesn't start new transaction because the implicit transaction has already started. The error above occurs because the snapshot was already created before the BEGIN command.

We can also see the similar error when using "BEGIN DEFERRABLE".

代码审查

- 审查流程

- 查找功能

- 代码审查

- 文档

- 功能

- 测试反馈

```
diff --git a/src/bin/pgbench/pgbench.c b/src/bin/pgbench/pgbench.c
index a3a3543..12a348fe 100644
--- a/src/bin/pgbench/pgbench.c
+++ b/src/bin/pgbench/pgbench.c
@@ -3070,6 +3070,39 @@ chooseScript(TState *thread)
     return i - 1;
 }

/* Prepare SQL commands in the chosen script */
static void
prepareCommands(CState *st)
{
+    int             j;
+    Command **commands = sql_script[st->use_file].commands;
+
+    for (j = 0; commands[j] != NULL; j++)
+    {
+        PQresult *res;
+        char          name[MAX_PREPARE_NAME];
+
+        if (commands[j]->type != SQL_COMMAND)
+            continue;
+        preparedStatementName(name, st->use_file, j);
+        res = PQprepare(st->con, name,
+                        commands[j]->args[0], commands[j]->args - 1, NULL);
+        if (!PQresultStatus(res) != PGRES_COMMAND_OK)
+            pg_log_error("is", PQerrorMessage(st->con));
+        PQclear(res);
+    }
+    st->prepared[st->use_file] = true;
+}
+
/* Send a SQL command, using the chosen querynode */
static bool
sendCommand(CState *st, Command *command)
@@ -3104,39 +3126,7 @@ sendCommand(CState *st, Command *command)
    const char *params[MAX_ARGS];
    if (!st->prepared[st->use_file])
    {
-        int             j;
-        Command **commands = sql_script[st->use_file].commands;
-
-        for (j = 0; commands[j] != NULL; j++)
-        {
-            PQresult *res;
-
-            if (commands[j]->type != SQL_COMMAND)
-                continue;
-            preparedStatementName(name, st->use_file, j);
-            if (PQpipelineStatus(st->con) == PQ_PIPELINE_OFF)
-            {
-                res = PQprepare(st->con, name,
-                               commands[j]->args[0], commands[j]->args - 1, NULL);
-                if (!PQresultStatus(res) != PGRES_COMMAND_OK)
-                    pg_log_error("is", PQerrorMessage(st->con));
-                PQclear(res);
-            }
-
```

代码审查

A brief SGML tutorial

- 审查流程

- 查找功能

- 代码审查

- 文档

- 功能

- 测试反馈

```
diff --git a/doc/src/sgml/ref/pg_waldump.sgml b/doc/src/sgml/ref/pg_waldump.sgml
index 5735a161ce..c953703bc8 100644
--- a/doc/src/sgml/ref/pg_waldump.sgml
+++ b/doc/src/sgml/ref/pg_waldump.sgml
@@ -100,6 +100,29 @@ PostgreSQL documentation
     </listitem>
     <varlistentry>
         <term><option>-k <replaceable>block</replaceable></option></term>
         <term><option>--block=<replaceable>block</replaceable></option></term>
     <listitem>
         <para>
             Display only records touching the given block. (Requires also
             providing the relation via <option>--relation</option>.)
         </para>
     </listitem>
     </varlistentry>

     <varlistentry>
         <term><option>-l <replaceable>tbl</replaceable>/<replaceable>db</replaceable>/<replaceable>rel</replaceable></option></term>
         <term><option>--relation=<replaceable>tbl</replaceable>/<replaceable>db</replaceable>/<replaceable>rel</replaceable></option></term>
     <listitem>
         <para>
             Display only records touching the given relation. The relation is
             specified via tablespace oid, database oid, and relfilenode separated
             by slashes.
         </para>
     </listitem>
     </varlistentry>

     <varlistentry>
         <term><option>-n <replaceable>limit</replaceable></option></term>
         <term><option>--limit=<replaceable>limit</replaceable></option></term>
```

Installation from Source Code

- 审查流程

- 查找功能

- 代码审查

- 文档

- 功能

- 测试反馈

```
$ git clone https://git.postgresql.org/git/postgresql.git
$ cd postgresql
$ git apply ../0001-Add-relation-block-filtering-to-pg_waldump.patch
$ mkdir build && cd build
$ ../configure --prefix=$PWD/pg --enable-tap-tests ...
$ make world && make check-world && make install-world
$ # 有时由于网络原因，构建文档的时候会失败
$ # 可以忽略文档的编译
$ # make install && (cd contrib; make install) && make check-world
$ cat <<END > pg-env.sh
export PGHOME=$PWD/pg
export PGDATA=\$PGHOME/pgdata
export PGPORT=5643
export PATH=\$PGHOME/bin:\$PATH
export LD_LIBRARY_PATH=\$PGHOME/lib:\$LD_LIBRARY_PATH
END
$ source pg-env.sh
$ initdb
$ pg_ctl -l log start
$ pg_waldump ...
```

Mailing List Etiquette FAQ

- 审查流程

- 查找功能

- 代码审查

- 文档

- 功能

- 测试反馈

```
02/25/2022 5      pgsql-hack Bharath Rupireddy      Re: [PATCH] add relation and block-level filtering to pg_waldump
02/25/2022 5      Japin Li          Re: [PATCH] add relation and block-level filtering to pg_waldump
-000:/*--F1 *mu4e-headers* 29% (614,0) (mu4e-headers Projectile EditorConfig (*) Undo-Tree yas company ivy WK) maildir
From: Japin Li <japinli@hotmail.com>
To: David Christensen <david.christensen@crunchydata.com>
Cc: pgsql-hackers@postgresql.org, pgsql-hackers@lists.postgresql.org
Subject: Re: [PATCH] add relation and block-level filtering to
pg_waldump
Flags: seen
Date: Fri Feb 25 21:08:20 2022
Maildir: /hotmail/Sent

On Fri, 25 Feb 2022 at 20:48, David Christensen <david.christensen@crunchydata.com> wrote:
>> Cool. I think we can report an error instead of reading wal files,
>> if the tablespace, database, or relation is invalid. Does there any
>> WAL record that has invalid tablespace, database, or relation OID?
>
> The only sort of validity check we could do here is range checking for the underlying data types
> (which we certainly could/should add if it's known to never be valid for the underlying types);

The invalid OID I said here is such as negative number and zero, for those
parameters, we do not need to read the WAL files, since it always invalid.

> non-existence of objects is a no-go, since that depends purely on the WAL range you are
> looking at and you'd have to, you know, scan it to see if it existed before marking as invalid. :)
>

Agreed.

-- 
Regards,
Japin Li.
```

代码贡献

- 编写代码 (Open Items)
- git diff & git format-patch



- 发送 Patch

Duplicate function call on timestamp2tm

From: U Japin <japini@outlook(dot)com>
To: "pgsql-hackers(at)postgresql(dot)org" <pgsql-hackers(at)postgresql(dot)org>
Subject: Duplicate function call on timestamp2tm
Date: 2019-12-12 14:34:20
Message-ID: 5G2PR06MB3776CA1450B96FCA7001EA9B6550B5G2PR06MB3776.apcrd06.prod.outlook.com
Views: Raw Message | Whole Thread | Download inbox | Resend email
Thread: 2019-12-12 14:34:20 from U Japin <japini@outlook(dot)com>
Lists: pgsql-hackers

Hi,
I find there is a duplicate function call on timestamp2tm in timestamp2tm_part and timestamp_part.
Is that necessary? I remove the latter one and it also works.

Best,
Japin.

Attachment	Content-Type	Size
remove-duplicate-timestamp2tm-function-call.patch	application/octet-stream	1.7 KB

代码贡献

① 创建 Patch

The screenshot shows a web interface for a Commitfest event. At the top, there are navigation links: 'Home' and 'Commitfest 2022-11'. On the right, it says 'Activity log / Logged in as japin (edit profile | log out)'. Below this, the title 'Commitfest 2022-11' is displayed. Underneath the title are three buttons: 'Search/filter', 'Shortcuts ▾', and 'New patch'. The 'New patch' button is highlighted with a blue border.

② 创建 Commitfest

The screenshot shows a 'New patch' form. At the top, there are navigation links: 'Home / Commitfest 2022-11 / New patch' and 'Logged in as japin (edit profile | log out)'. The main title is 'New patch'. Below the title are several input fields and labels: 'Description:' with a large text input field, 'Topic:' with a dropdown menu containing a single item '-----', 'Specify thread msgid:' with a text input field and a 'Find thread' button, and a 'Create patch' button at the bottom.

③ 审查 & 修改

代码贡献

PostgreSQL Patch Tester

[Current commitfest](#) | [Next commitfest](#) | [FAQ](#)

Here lives an experimental bot that converts email threads that are registered in the [Commitfest system](#) into [branches on Github](#), and collates test results from [Cirrus CI](#).

Ready for Committer

① 创建 Patch

40/3464	Add Amcheck option for checking unique constraints in btree indexes	Anastasia Lubennikova, Maxim Orlov, Pavel Borisov	≡ H						
40/3405	Add connection active_idle time to pg_stat_activity	Rafia Sabih, Sergey Dudoladov	≡ H						
40/3895	Avoid to handle FORCE_NOT_NULL/FORCE_NULL options when COPY TO	Mingli Zhang	≡ H						
40/2601	Fast COPY FROM command for the foreign tables	Andrey Lepikhov	≡ H						
40/2897	Faster pglz compression	Andrey Borodin, Владимир Лесков	≡ H						
40/3662	Fix assertion failure with barriers in parallel hash join	Thomas Munro, Melanie Plageman	≡ H						
40/3779	Improve description of XLOG_RUNNING_XACTS	Masahiko Sawada	≡ H						
40/3271	KnownAssignedXidsGetAndSetXmin performance	Michail Nikolaev	≡ H						
40/2903	Parallel Hash Full Join	Melanie Plageman	≡ H						

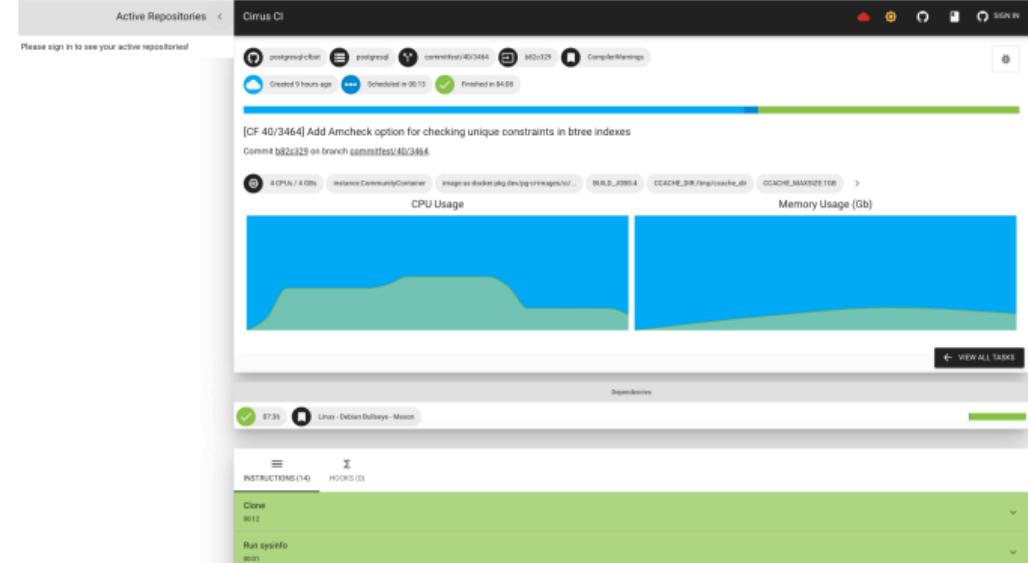
② 创建 Commitfest

③ 审查 & 修改

<http://cfbot.cputube.org/>

代码贡献

① 创建 Patch



② 创建 Commitfest

③ 审查 & 修改

代码贡献

① 创建 Patch

From: Tom Lane <tgl(at)sss(dot)pgh(dot)pa(dot)us>
To: Li Japin <japinli(at)hotmail(dot)com>
Cc: "pgsql-hacker(at)postgresql(dot)org" <pgsql-hackers(at)postgresql(dot)org>
Subject: Re: Duplicate function call on timestamp2tm
Date: 2019-12-12 14:41:18
Message-ID: 4024.1576161678@sss.pgh.pa.us
Views: Raw Message | Whole Thread | Download mbox | Resend email
Lists: pgsql-hackers

Li Japin <japinli(at)hotmail(dot)com> writes:
> I find there is a duplicate function call on timestamp2tm in timestamptz_part and timestamp_part.
> Is that necessary? I remove the latter one and it also works.

Huh. I do believe you're right. Must be an ancient copy-and-paste mistake?

regards, tom lane

② 创建 Commitfest

From: Tom Lane <tgl(at)sss(dot)pgh(dot)pa(dot)us>
To: Li Japin <japinli(at)hotmail(dot)com>
Cc: Greg Stark <stark(at)mit(dot)edu>, "pgsql-hacker(at)postgresql(dot)org" <pgsql-hackers(at)postgresql(dot)org>
Subject: Re: Duplicate function call on timestamp2tm
Date: 2019-12-12 15:13:17
Message-ID: 5511.1576163597@sss.pgh.pa.us
Views: Raw Message | Whole Thread | Download mbox | Resend email
Lists: pgsql-hackers

I wrote:
> Li Japin <japinli(at)hotmail(dot)com> writes:
>> I find there is a duplicate function call on timestamp2tm in timestamptz_part and timestamp_part.
>> Is that necessary? I remove the latter one and it also works.

> Huh. I do believe you're right. Must be an ancient copy-and-paste mistake?

③ 审查 & 修改

Ah, after looking in the git history, not quite that ancient: this duplication dates to commit 258ee1b63, which moved these switch cases from the "if (type == RESERV)" switches in the same functions. In the previous coding these function calls were actually necessary, but here they're redundant. I guess that's just additional ammunition for Greg's point that the keywords were misclassified ;-).

I see from the code coverage report that we're missing coverage for these and some other paths in timestamp[tz].part. Think I'll go add some more test cases while I'm at it.

Thanks again for the report!

资源汇总

<http://cfbot.cputube.org/>
<https://babel.postgresql.org>
<https://buildfarm.postgresql.org>
<https://commitfest.postgresql.org/>
https://wiki.postgresql.org/wiki/Development_information
https://wiki.postgresql.org/wiki/PostgreSQL_15_Open_Items
https://wiki.postgresql.org/wiki/PostgreSQL_Related_Slides_and_Presentations
<https://www.postgresql.org/account/submitbug/>
<https://www.postgresql.org/list/>

